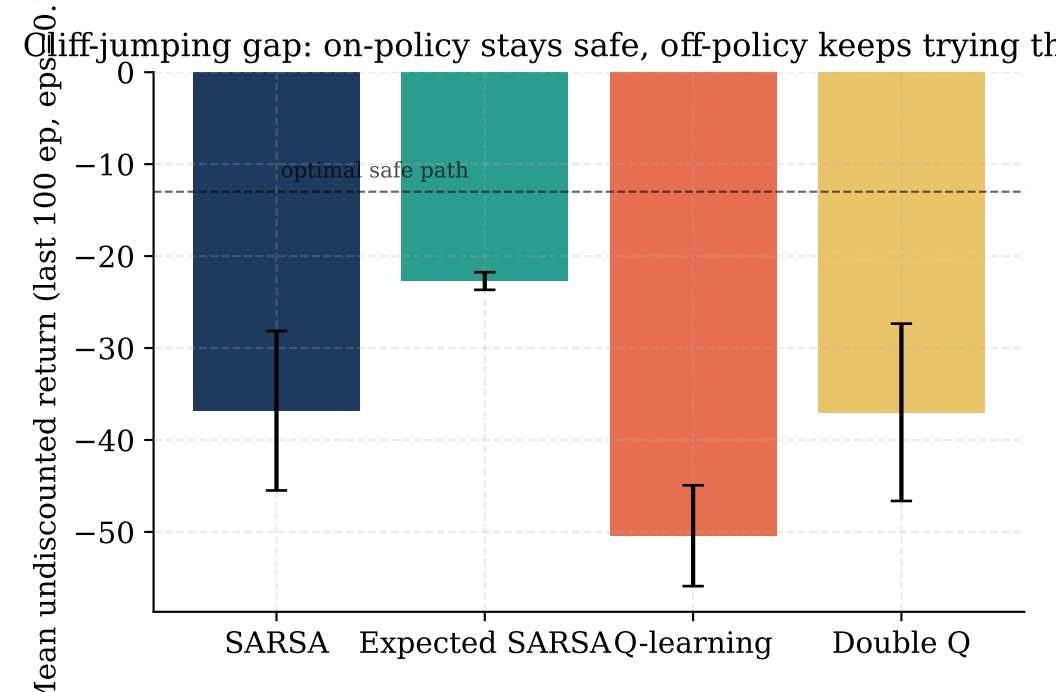
Cliff-jumping gap: on-policy stays safe, off-policy keeps trying th
Mean undiscounted return (last 100 ep, eps=0.
0
−10
−20
−30
−40
−50
optimal safe path
SARSA
Expected SARSA
Q-learning
Double Q