Week 10
SAC, TD3, and CrossQ for Continuous Control

STUDY-RL — a self-paced Deep RL curriculum

November 8, 2026

Contents

Learning objectives 3

1 Motivation and intuition 4

2 Mathematical formulation 5
- 2.1 Stochastic vs. deterministic policies . . . 5
- 2.2 Maximum entropy RL . . . 5
- 2.3 Tanh-squashed Gaussian + Jacobian correction . . . 6
- 2.4 Automatic temperature tuning . . . 6
- 2.5 TD3 . . . 6
- 2.6 CrossQ and SimBa: distinct design choices . . . 7

3 Algorithms 8

4 Implementation walkthrough 9
- 4.1 Tanh-squashed Gaussian (the right way) . . . 9
- 4.2 SAC critic update . . . 9

5 Empirical results 10
- 5.1 Statistics bridge: The minimum of two unbiased critics can be biased . . . 10

6 Safety angle 12

7 Matching application 13

8 Exercises 14

9 Further reading 15

Learning objectives

Study companion. Read `teaching_guide.md` in this folder for step-by-step examples, an analogy with its limits, and self-checks. Then use `workbook.ipynb` to implement and vary the calculations.

Key Idea

By the end of this week you will be able to:

1. Contrast stochastic and deterministic policy families and identify when each yields a lower-variance gradient.
2. Derive the SAC (soft actor-critic) objective from the maximum-entropy MDP formulation; state the soft Bellman equations and the closed-form soft policy.
3. Implement the tanh-squashed Gaussian policy, including the Jacobian correction to its log-density, and explain why the correction is essential.
4. Train SAC on Pendulum-v1 (or MuJoCo Hopper) with automatic temperature tuning; plot return and entropy.
5. Specify TD3 (twin critics, target policy smoothing, delayed updates) and explain its role as the deterministic-policy counterpart to SAC.
6. Summarise CrossQ (2024) and the SimBa block (2025) as published architecture and training alternatives, including why CrossQ drops target networks.

1 Motivation and intuition

PPO is on-policy: it reuses each rollout for several optimisation epochs, then normally replaces it with fresh data. For continuous control with expensive simulators (or robots), this is a luxury.

Off-policy methods — SAC, TD3, and the modern ensemble variants — amortise every transition across many gradient updates through a replay buffer, and they exploit the reparameterisation trick (week 7) to compute low-variance pathwise gradients for the actor.

Replay ratio counts gradient updates per environment step. It is not the expected number of uses of each transition; batch size and buffer residence also matter.

Intuition. SAC and TD3 are duals.

- TD3 trains a *deterministic* policy $a = \mu_\theta(s)$. Exploration comes from Gaussian noise added to actions at *rollout* time, not from the policy itself. Stability tricks: twin critics, delayed actor updates, target-action smoothing.
- SAC trains a *stochastic* policy $\pi_\theta(a \mid s) = \tanh(\mathcal{N}(\mu_\theta(s), \sigma_\theta(s)^2))$. Exploration is a property of the policy itself. The critic learns the soft Q-function; the actor maximises a soft-Bellman-corrected entropy-regularised objective.

Both train two critics and use target networks; both update the actor by gradient descent through the critic (pathwise gradient).

Common pitfall

The single most common bug in a SAC implementation is forgetting the tanh-Jacobian correction in the log-probability. The policy will then silently collapse to a near-deterministic edge of the action space, scoring artificially high entropy reported by the un-corrected formula. Verify the density, not its logarithm: $\int_{-1}^{1} \pi(a \mid s)\, da = 1$. For one action dimension and $U \sim \text{Uniform}(-1, 1)$, $2\mathbb{E}[\exp(\log \pi(U \mid s))] = 1$; numerical quadrature is another check.

2 Mathematical formulation

2.1 Stochastic vs. deterministic policies

Stochastic policy gradient (score-function form): $\nabla_\theta J = \mathbb{E}[\nabla_\theta \log \pi_\theta(a \mid s)\, Q^\pi(s,a)]$.

Deterministic policy gradient (Silver 2014):

Theorem 2.1.1 (DPG). *For a deterministic policy $a = \mu_\theta(s)$ and a state distribution d^μ,*

$$\nabla_\theta J(\mu_\theta) = \mathbb{E}_{s\sim d^\mu}\left[\nabla_\theta \mu_\theta(s)\, \nabla_a Q^{\mu_\theta}(s,a)\big|_{a=\mu_\theta(s)}\right].$$

Sketch. Start from the start-state return $J = \mathbb{E}_{s_0\sim\rho_0} V^\mu(s_0)$, differentiate its Bellman recursion, and unroll the future-state derivative. The resulting discounted occupancy weights the direct actor–critic chain rule. If $d^\mu = (1-\gamma)\sum_t \gamma^t P(S_t = \cdot)$ is normalised, the displayed expectation needs a factor $1/(1-\gamma)$; it is exact as written with the unnormalised occupancy measure. Replay states instead produce a practical off-policy surrogate, not automatically the same gradient. See https://proceedings.mlr.press/v32/silver14.html. □

Deterministic policy gradients require action differentiability and remove explicit action sampling in this expression. There is no universal variance ordering against every stochastic-policy estimator.

2.2 Maximum entropy RL

The classical RL objective $J(\theta) = \mathbb{E}[\sum_t \gamma^t r_t]$ is augmented to

$$J_{\text{MaxEnt}}(\theta) = \mathbb{E}\left[\sum_t \gamma^t (r_t + \alpha \mathcal{H}(\pi_\theta(\cdot \mid s_t)))\right],$$

with $\alpha \geq 0$ the temperature. The optimal policy is no longer deterministic; it is the Gibbs distribution over actions weighted by Q.

Theorem 2.2.1 (Soft Bellman equations, Haarnoja 2018a). *Assume $\alpha > 0$, bounded rewards, $\gamma < 1$, and a finite partition integral relative to the chosen action measure. For this entropy-regularised MDP, the soft value function $V^\pi(s) = \mathbb{E}_{a\sim\pi}[Q^\pi(s,a) - \alpha \log \pi(a \mid s)]$, and the soft Bellman backup*

$$\mathcal{T}^\pi Q(s,a) \;=\; r(s,a) + \gamma\, \mathbb{E}_{s'}\left[\mathbb{E}_{a'\sim\pi}[Q(s',a') - \alpha \log \pi(a' \mid s')]\right].$$

The soft-optimal policy is the Boltzmann form $\pi^\star(a \mid s) \propto \exp(Q^\star(s,a)/\alpha)$.

Derivation of the optimal policy. At fixed Q, the actor update solves

$$\pi^{\text{new}}(\cdot \mid s) = \arg\max_\pi \mathbb{E}_{a\sim\pi}[Q(s,a)] + \alpha \mathcal{H}(\pi(\cdot \mid s)).$$

Lagrangian on the simplex $\int \pi(a \mid s)\, da = 1$:

$$\mathcal{L}(\pi, \nu) = \int \pi(a \mid s)\, Q(s,a)\, da - \alpha \int \pi \log \pi\, da - \nu\Big(\int \pi\, da - 1\Big).$$

Functional derivative: $\delta\mathcal{L}/\delta\pi = Q(s,a) - \alpha(\log \pi + 1) - \nu = 0$. Solving, $\pi(a \mid s) = \exp((Q(s,a) - \nu - \alpha)/\alpha)$, and normalising gives the Boltzmann form. The result $V = \alpha \log \int \exp(Q/\alpha)\, da$ is the soft-max (log-sum-exp) over actions. □

2.3 Tanh-squashed Gaussian + Jacobian correction

We need a stochastic policy that:

1. samples d_a action coordinates in a bounded box $[-1, 1]^{d_a}$ (continuous-control envs require this);
2. is differentiable through the sample (reparameterisable);
3. has a tractable density.

The standard recipe: $u \sim \mathcal{N}(\mu_\theta(s), \sigma_\theta(s)^2)$, then $a = \tanh(u)$. By the change-of-variables formula,

$$\pi(a \mid s) = \mathcal{N}(u; \mu, \sigma^2) \cdot \left| \frac{du}{da} \right|.$$

Since $u = \tanh^{-1}(a)$, $du/da = 1/(1 - a^2)$, while $da/du = 1 - \tanh^2(u)$. Taking logarithms gives

$$\log \pi(a \mid s) = \log \mathcal{N}(u; \mu, \sigma^2) - \sum_{i=1}^{d_a} \log(1 - a_i^2).$$

Numerically, $\log(1 - \tanh^2 u) = 2(\log 2 - u - \mathrm{softplus}(-2u))$ is stable.

Common pitfall

Drop the Jacobian and your policy will report $\log \pi$ that is artificially *lower* (because tanh compresses); the entropy bonus $-\alpha \log \pi$ in the actor loss becomes artificially high; the policy will push to the boundary of the action space and saturate. The optimiser will not crash; it will silently converge to a wrong policy.

2.4 Automatic temperature tuning

Haarnoja 2018b makes α a learned parameter by enforcing the constraint

$$\mathbb{E}_{a \sim \pi_\theta}[-\log \pi_\theta(a \mid s)] \geq \mathcal{H}_{\text{target}} = -d_a.$$

The dual update is

$$\alpha \leftarrow \alpha - \eta_\alpha \nabla_\alpha \alpha \cdot (-\log \pi_\theta(a \mid s) - \mathcal{H}_{\text{target}}),$$

implemented by storing $\log \alpha$ unconstrained and computing $\alpha = \exp(\log \alpha)$. This adapts temperature but retains a target-entropy choice, optimiser settings, and action-coordinate convention. It is not tuning-free.

2.5 TD3

Twin Delayed DDPG (Fujimoto 2018) fixes three pathologies of DDPG:

1. **Overestimation bias.** Train two critics Q_{ϕ_1}, Q_{ϕ_2}; the target uses $\min(Q_{\phi'_1}, Q_{\phi'_2})$.
2. **Function-approximation error in actor target.** Delay actor updates: update the actor every $d = 2$ critic steps.
3. **Deterministic-policy target sharpness.** Smooth the target action: $a' = \mu_{\theta'}(s') + \mathrm{clip}(\epsilon, -c, c)$, $\epsilon \sim \mathcal{N}(0, \tilde{\sigma}^2)$.

The TD3 target is

$$y = r + \gamma(1 - d) \min_{i=1,2} Q_{\phi'_i}(s', a'), \quad a' = \mathrm{clip}(\mu_{\theta'}(s') + \mathrm{clip}(\epsilon, -c, c), a_{\min}, a_{\max}).$$

2.6 CrossQ and SimBa: distinct design choices

CrossQ uses batch normalisation with current and successor critic inputs processed together, while stopping gradients through the target and omitting target networks. Its experiments use a low update-to-data ratio; this does not establish that arbitrary BatchNorm placement stabilises TD learning. See https://arxiv.org/abs/1902.05605. SimBa is a separate scalable architecture: running observation normalisation, residual feedforward blocks and layer normalisation. RMS normalisation and resets are not a definition of the original SimBa architecture. See https://arxiv.org/abs/2410.09754.

3 Algorithms

SAC with automatic temperature

Hyper-params: replay buffer D of size 10^6, batch $B = 256$, $\gamma = 0.99$, polyak $\tau = 5 \times 10^{-3}$, actor/critic lr $= 3 \times 10^{-4}$, α-lr $= 3 \times 10^{-4}$, $\mathcal{H}_{\text{target}} = -d_a$.
Init: actor π_θ, critics Q_{ϕ_1}, Q_{ϕ_2} and their polyak copies $Q_{\phi'_1}, Q_{\phi'_2}$.
For each env step:

1. Sample $a \sim \pi_\theta(\cdot \mid s)$, step env, store (s, a, r, s', d).
2. Sample minibatch $\sim D$.
3. Compute target: $\tilde{a} \sim \pi_\theta(\cdot \mid s')$, $\log \tilde{p} = \log \pi_\theta(\tilde{a} \mid s')$,
$$y = r + \gamma(1 - d)\big[\min_i Q_{\phi'_i}(s', \tilde{a}) - \alpha \log \tilde{p}\big].$$
4. Critic update: minimise $\sum_i (Q_{\phi_i}(s, a) - y)^2$.
5. Actor update (reparameterised): minimise $\mathbb{E}[\alpha \log \pi_\theta(a \mid s) - \min_i Q_{\phi_i}(s, a)]$.
6. Temperature update: minimise $\mathbb{E}[-\alpha(\log \pi_\theta(a \mid s) + \mathcal{H}_{\text{target}})]$ in $\log \alpha$.
7. Polyak update: $\phi'_i \leftarrow \tau\phi_i + (1 - \tau)\phi'_i$.

TD3

Hyper-params: as SAC, plus action-noise $\sigma = 0.1$, target-noise $\tilde{\sigma} = 0.2$, noise clip $c = 0.5$, delay $d = 2$.
For each env step:

1. Roll out with $a = \mu_\theta(s) + \mathcal{N}(0, \sigma^2)$.
2. Critic step (twin): use the terminal-masked target above, with noisy next actions clipped to the environment bounds.
3. Every d steps: actor step ($\nabla_\theta \mathbb{E}[Q_{\phi_1}(s, \mu_\theta(s))]$), then polyak update of all targets.

4 Implementation walkthrough

4.1 Tanh-squashed Gaussian (the right way)

```
class TanhGaussianPolicy(nn.Module):
    def __init__(self, obs_dim, act_dim, hidden=(256, 256)):
        super().__init__()
        # mu and log_sigma heads share a trunk; both depend on s.
        self.trunk = mlp([obs_dim, *hidden], output_activation=nn.ReLU)
        self.mu = nn.Linear(hidden[-1], act_dim)
        self.log_sigma = nn.Linear(hidden[-1], act_dim)

    def sample(self, obs):
        z = self.trunk(obs)
        mu = self.mu(z)
        log_sigma = self.log_sigma(z).clamp(-5, 2)
        sigma = log_sigma.exp()
        eps = torch.randn_like(mu)
        u = mu + sigma * eps
        a = torch.tanh(u)
        # log N(u; mu, sigma^2) - sum log(1 - a^2)  (numerically stable form)
        log_p_u = -0.5 * (((u - mu) / sigma) ** 2 + 2 * log_sigma
                          + math.log(2 * math.pi))
        # log|da/du|^-1 = -log(1 - tanh^2 u); stable form:
        log_det_jac = 2.0 * (math.log(2.0) - u - F.softplus(-2.0 * u))
        log_prob = (log_p_u - log_det_jac).sum(-1)
        return a, log_prob
```

4.2 SAC critic update

```
def critic_update(batch, actor, q1, q2, q1_tgt, q2_tgt, log_alpha,
                  gamma=0.99, opt=None):
    s, a, r, s2, d = batch
    with torch.no_grad():
        a2, logp2 = actor.sample(s2)
        q_tgt = torch.min(q1_tgt(s2, a2), q2_tgt(s2, a2)) - log_alpha.exp() * logp2
        y = r + gamma * (1 - d) * q_tgt
    L = ((q1(s, a) - y) ** 2 + (q2(s, a) - y) ** 2).mean()
    opt.zero_grad(); L.backward(); opt.step()
```

5 Empirical results

5.1 Statistics bridge: The minimum of two unbiased critics can be biased

Example 5.1.1 (Original statistical teaching fixture)**.** Both critics estimate a true scalar value 5 with independent errors +1 or -1, each equally likely. This artificial error model isolates the minimum operation.

The four equally likely critic pairs are $(4,4),(4,6),(6,4),(6,6)$. Their minima are $4,4,4,6$, with mean $18/4 = 4.5$. Each individual critic has mean 5, but their minimum has bias -0.5.

Clipped double estimates can counter overestimation without being generically unbiased or guaranteeing a better final policy. Critic errors are often correlated because they share replay data and targets; the fixture explicitly assumed independence.

Check a changed assumption. Make both errors identical on every draw. What is the expected minimum? Does this show that a low empirical critic disagreement certifies accuracy?

Check. The minimum is then the shared critic estimate, with mean 5 in this fixture. Both critics could instead share a large systematic bias, so agreement alone does not certify accuracy.

Statistical foundation. Siegrist § § 4.1, 7.1, PDF pp.439, 896. Section source: https://stats.libretexts.org/@go/page/10189. The course statistics companion records the supplied editions and one-based PDF pages. This fixture is not a newly measured project result.

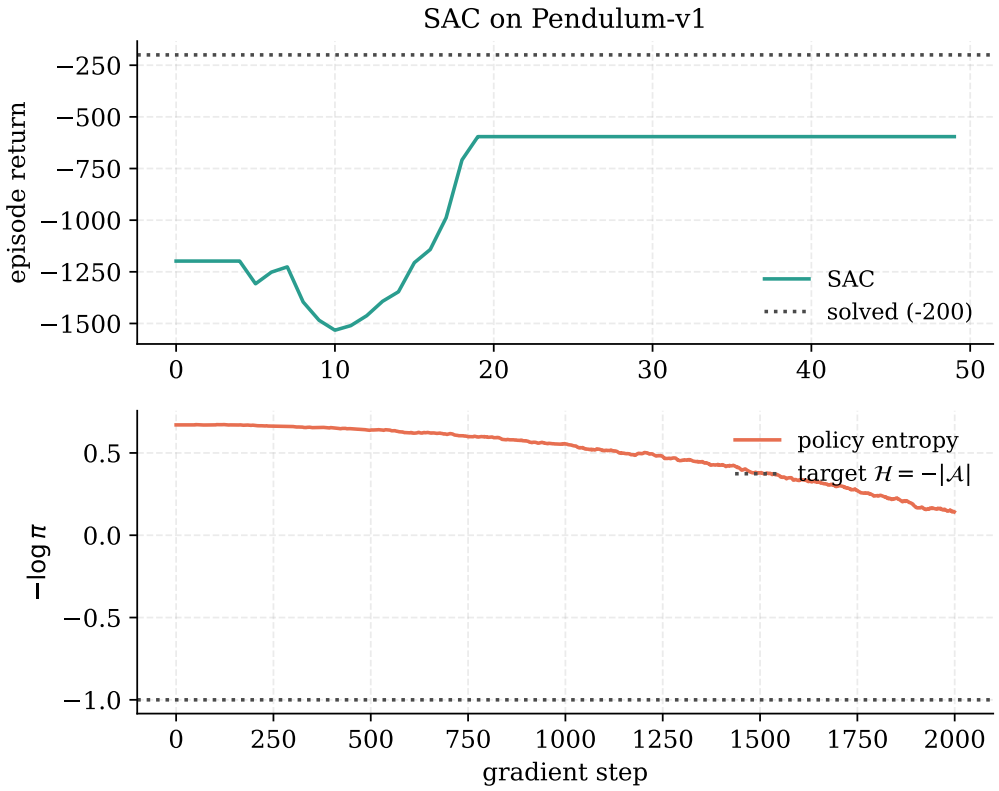


Figure 5.1: SAC on Pendulum-v1. Three seeds; return (top) and policy entropy (bottom). Note that the entropy decreases during training as the policy becomes more confident; with automatic temperature tuning, it stabilises near the target $-d_a = -1$.

6 Safety angle

Safety angle

Three failure modes that matter for high-stakes off-policy RL.

1. **Q-function over-estimation cascade.** A single mis-fit batch can push $Q(s, a)$ several units too high; the next policy update will increase the probability of that action, generating more high-Q data, generating more upward bias. The twin-critic min trick of TD3/SAC is the structural fix. *Watch the empirical* $\min_i Q_i$ *during training*; a runaway uptrend is the failure signature.

2. **Stuck entropy.** If the automatic temperature update is mis-coded (most commonly: mixing the entropy-target sign or differentiating through the actor during a temperature step), the policy entropy either explodes or collapses. Mitigation: clamp $\log \alpha \in [-5, 2]$ and assert the gradient sign is correct in unit tests.

3. **Replay-buffer drift.** On long runs ($> 10^6$ steps) the buffer contains transitions from radically different past policies. The Q-targets become biased. This is exactly the setting where *constrained safe RL* (week 22) needs explicit safety filters; the off-policy bias compounds with reward mis-specification to produce non-monotone training.

Matching context. For an off-policy matching system retrained nightly on the day's logs, the entropy-regularised objective is structurally useful: it prevents the policy from collapsing onto a single broker / candidate / item even when the reward signal degenerates. The same maximum-entropy principle is the rationale for the KL-to-reference term in RLHF (week 18); the only difference is that there the "reference" is a fixed prior policy, while in SAC it is the uniform distribution.

7 Matching application

Matching application

SAC as the nightly retrainer for synthetic ATS candidate-ranking.
The matching setting in production usually has continuous-valued match-quality features and is too high-dimensional for tabular methods. SAC fits naturally:

- **State s:** candidate embedding (768-dim), job-pool snapshot (256-dim).
- **Action $a \in \mathbb{R}^d$:** score adjustments over the top-100 candidates (so $d_a = 100$, then we softmax-sample the top-K).
- **Reward r:** hire-conversion within 30 days, possibly normalised by base rate.
- **Replay buffer D:** 90 days of logged (s, a, r, s').

The automatic temperature tuning is particularly valuable here: it adapts the temperature around a chosen entropy target, while that target and learning settings still influence exploration. The maximum-entropy objective keeps the candidate-shortlisting policy diverse, which encourages diversity but provides no individual or group fairness guarantee; an entropy bonus can coexist with severe exposure disparities.

Operational caveats.

- Production deployment requires a CMDP formulation (week 22) for expected-cost constraints (e.g. minimum exposure per demographic group). SAC alone is the optimiser; it does not provide guarantees.
- Replay-buffer drift over the 90-day window is a real concern; in practice we re-fit the value head weekly from scratch on the most recent two weeks to combat this.

8 Exercises

E1 ($\star$) Verify the tanh-Jacobian correction: sample $u \sim \mathcal{N}(0,1)$, transform to $a = \tanh u$, compute the corrected $\log \pi(a)$, and check that the density $\exp(\log \pi(a))$ integrates to 1 over $(-1, 1)$.

E2 ($\star$) Derive the soft-optimal policy in closed form via the Lagrangian and confirm $\pi^\star \propto \exp(Q/\alpha)$.

E3 ($\star\star$) Implement the SAC critic update for Pendulum and verify the loss decreases over 200 steps.

E4 ($\star\star$) Implement the actor reparameterised loss; verify the gradient sign on a toy 1-d problem.

E5 ($\star\star$) Implement automatic temperature tuning; show that α converges to a value that yields $\mathcal{H} \approx \mathcal{H}_{\text{target}}$.

E6 ($\star\star$) Train SAC on Pendulum-v1 to a smoothed return ≥ -200 within 8000 env steps.

E7 ($\star\star\star$) Code TD3 in $\sim$100 lines; compare wall-clock to SAC on Pendulum.

E8 ($\star\star\star$, meaningful objective) For a logged-bandit matching dataset, run an off-policy SAC variant: load the logs into the replay buffer and train without env interaction. compare held-out estimates, support and uncertainty with the behaviour policy. Vanilla offline SAC can extrapolate unsupported actions; an apparent SNIPS gain on the training log is not a reliable improvement certificate.

9 Further reading

See `extended_reading.pdf`. Five papers: Lillicrap et al. 2015 (DDPG), Fujimoto et al. 2018 (TD3), Haarnoja et al. 2018a/b (SAC + automatic temperature), Bhatt et al. 2024 (CrossQ), Lee et al. 2025 (SimBa). Sourced from `research/deeprl_core_sota.md` §3.1–3.7.