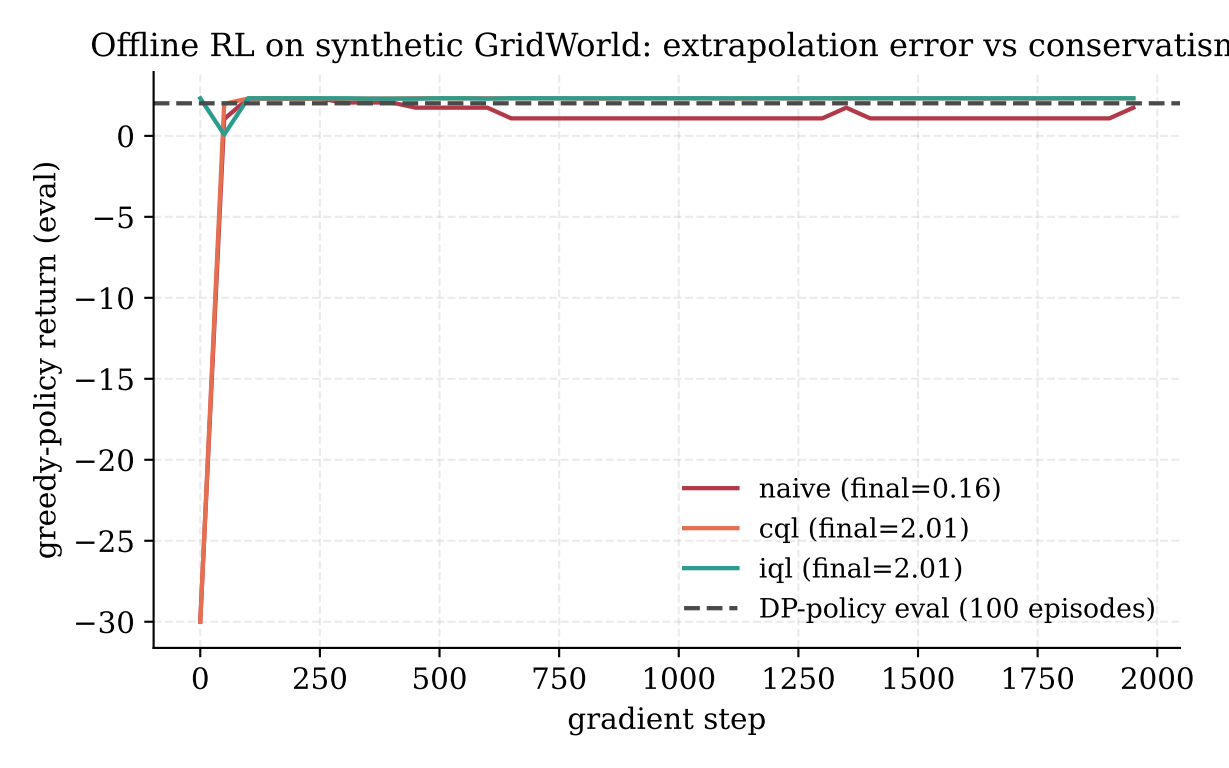
Offline RL on synthetic GridWorld: extrapolation error vs conservatism
greedy-policy return (eval)
0
−5
−10
−15
−20
−25
−30
0
250
500
750
1000
1250
1500
1750
2000
gradient step
naive (final=0.16)
cql (final=2.01)
iql (final=2.01)
DP-policy eval (100 episodes)