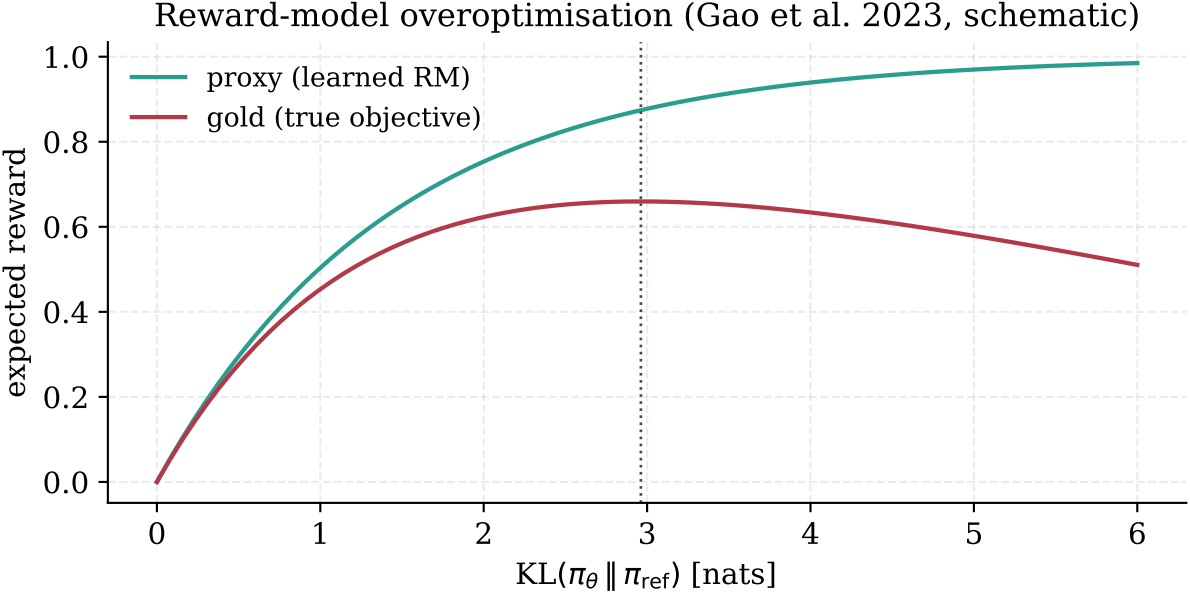
Reward-model overoptimisation (Gao et al. 2023, schematic)
proxy (learned RM)
gold (true objective)
expected reward
1.0
0.8
0.6
0.4
0.2
0.0
0
1
2
3
4
5
6
KL($\pi_\theta \,\|\, \pi_{\text{ref}}$) [nats]